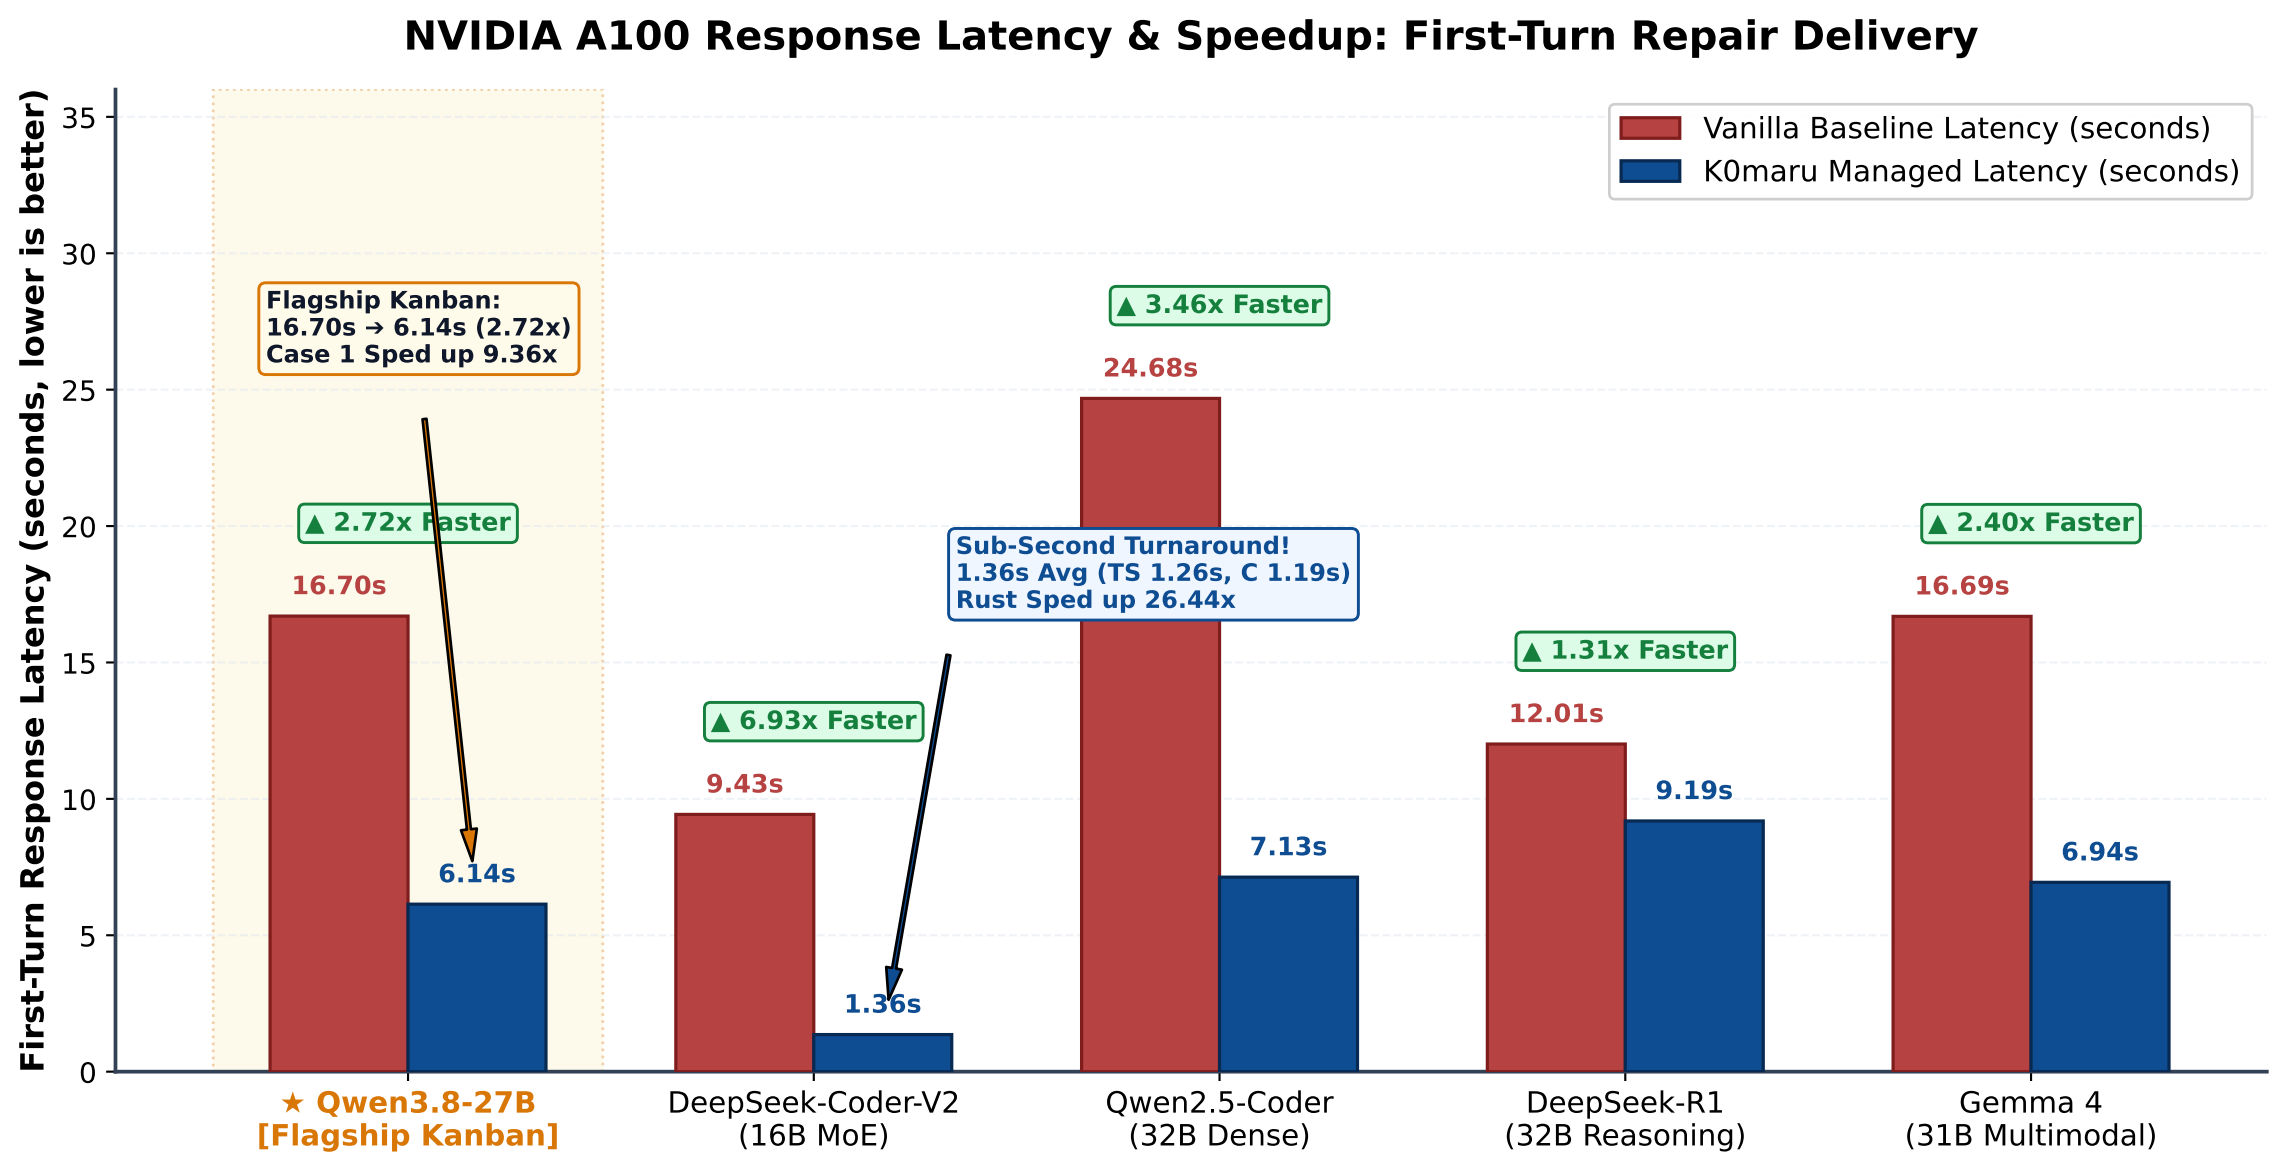
NVIDIA A100 Response Latency & Speedup: First-Turn Repair Delivery
First-Turn Response Latency (seconds, lower is better)
0
5
10
15
20
25
30
35
Vanilla Baseline Latency (seconds)
K0maru Managed Latency (seconds)
Flagship Kanban:
16.70s → 6.14s (2.72x)
Case 1 Sped up 9.36x
▲ 2.72x Faster
16.70s
6.14s
▲ 6.93x Faster
9.43s
1.36s
Sub-Second Turnaround!
1.36s Avg (TS 1.26s, C 1.19s)
Rust Sped up 26.44x
▲ 3.46x Faster
24.68s
7.13s
▲ 1.31x Faster
12.01s
9.19s
▲ 2.40x Faster
16.69s
6.94s
★ Qwen3.8-27B
[Flagship Kanban]
DeepSeek-Coder-V2
(16B MoE)
Qwen2.5-Coder
(32B Dense)
DeepSeek-R1
(32B Reasoning)
Gemma 4
(31B Multimodal)