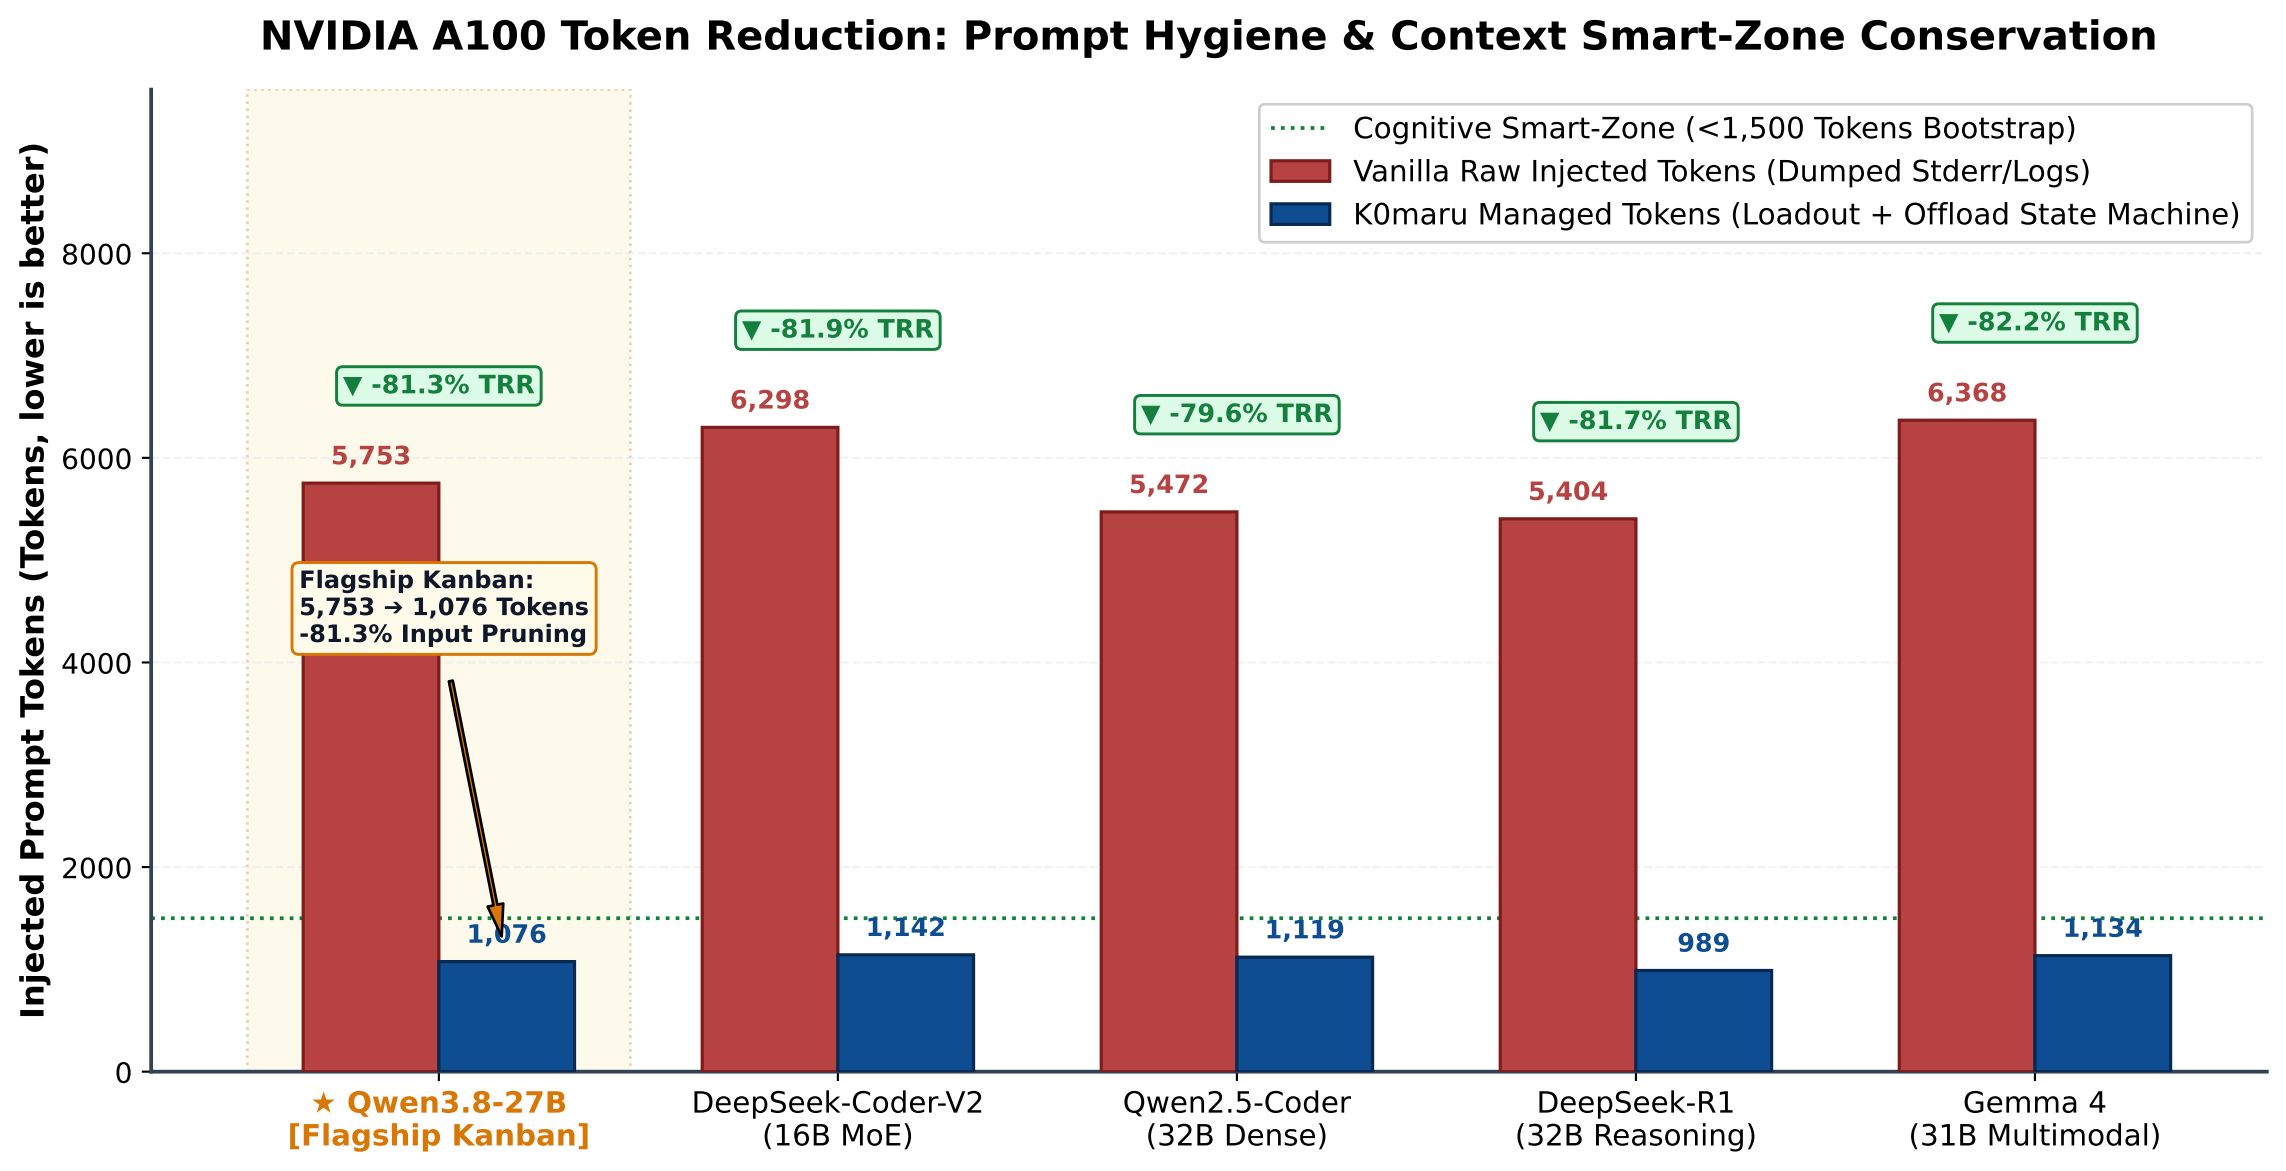
NVIDIA A100 Token Reduction: Prompt Hygiene & Context Smart-Zone Conservation
Injected Prompt Tokens (Tokens, lower is better)
Cognitive Smart-Zone (<1,500 Tokens Bootstrap)
Vanilla Raw Injected Tokens (Dumped Stderr/Logs)
K0maru Managed Tokens (Loadout + Offload State Machine)
8000
6000
4000
2000
0
▼ -81.3% TRR
5,753
1,076
Flagship Kanban:
5,753 → 1,076 Tokens
-81.3% Input Pruning
▼ -81.9% TRR
6,298
1,142
▼ -79.6% TRR
5,472
1,119
▼ -81.7% TRR
5,404
989
▼ -82.2% TRR
6,368
1,134
★ Qwen3.8-27B
[Flagship Kanban]
DeepSeek-Coder-V2
(16B MoE)
Qwen2.5-Coder
(32B Dense)
DeepSeek-R1
(32B Reasoning)
Gemma 4
(31B Multimodal)